

Imitative Planning for Autonomous Vehicles



Nicholas Rhinehart
Carnegie Mellon University



Rowan McAllister
UC Berkeley



Kris Kitani
Carnegie Mellon University



Sergey Levine
UC Berkeley



Q: Can robots **safely learn** to drive **suburban roads** in **interpretable** ways to **new goals** ?



Q: Can robots **safely learn** to drive **suburban roads** in **interpretable** ways to **new goals** ?

Without learning ?

Behavior cloning ?

Model-based RL ?

Q: Can robots **safely learn** to drive **suburban roads** in **interpretable** ways to **new goals** ?
(trains offline)

Without learning ?

Behavior cloning ?

Model-based RL ?

Q: Can robots **safely learn** to drive **suburban roads** in **interpretable** ways to **new goals** ?
(trains offline)

Without learning ?



Behavior cloning ?



Model-based RL ?



Q: Can robots **safely learn** to drive **suburban roads** in **interpretable** ways to **new goals** ?

(trains offline)

(no rewards required)

Without learning ?



Behavior cloning ?



Model-based RL ?



Q: Can robots **safely learn** to drive **suburban roads** in **interpretable** ways to **new goals** ?

(trains offline)

(no rewards required)

Without learning ?



Behavior cloning ?



Model-based RL ?



Q: Can robots **safely learn** to drive **suburban roads** in **interpretable** ways to **new goals** ?

(trains offline)

(no rewards required)

(plans)

| | | |
|--------------------|---|---|
| Without learning ? |  |  |
| Behavior cloning ? |  |  |
| Model-based RL ? |  |  |

Q: Can robots **safely learn** to drive **suburban roads** in **interpretable** ways to **new goals** ?

(trains offline)

(no rewards required)

(plans)

Without learning ?



Behavior cloning ?



Model-based RL ?



Q: Can robots **safely learn** to drive **suburban roads** in **interpretable** ways to **new goals** ?

(trains offline)

(no rewards required)

(plans)

(dyn. model)

Without learning ?



Behavior cloning ?



Model-based RL ?




Q: Can robots **safely learn** to drive **suburban roads** in **interpretable** ways to **new goals** ?

| | (trains offline) | (no rewards required) | (plans) | (dyn. model) |
|--------------------|------------------|-----------------------|---------|--------------|
| Without learning ? | ✓ | ✗ | ✓ | ? |
| Behavior cloning ? | ? | ✓ | ✗ | ✗ |
| Model-based RL ? | ? | ✗ | ✓ | ✓ |



Q: Can robots **safely learn** to drive **suburban roads** in **interpretable** ways to **new goals** ?

| | (trains offline) | (no rewards required) | (plans) | (dyn. model) |
|-------------------------|------------------|-----------------------|---------|--------------|
| Without learning ? | ✓ | ✗ | ✓ | ? |
| Behavior cloning ? | ? | ✓ | ✗ | ✗ |
| Model-based RL ? | ? | ✗ | ✓ | ✓ |
| Imitative Models (ours) | ✓ | ✓ | ✓ | ✓ |



Modelling Expert Drivers



Modelling Expert Drivers





Modelling Expert Drivers





$$\text{states} = f(\text{latents}; \text{context})$$

Modelling Expert Drivers



$$\begin{aligned}\text{states} &= f(\text{latents}; \text{context}) \\ \text{latents} &= f^{-1}(\text{states}; \text{context})\end{aligned}$$

Modelling Expert Drivers





$$\text{states} = f(\text{latents}; \text{context})$$

$$\text{latents} = f^{-1}(\text{states}; \text{context})$$

$$q(\text{states} \mid \text{context}) = \frac{\mathcal{N}(\text{latents}; 0, I)}{\left| \det \frac{\partial f}{\partial \text{latents}} \right|}$$

Modelling Expert Drivers




$$\begin{aligned}\text{states} &= f(\text{latents}; \text{context}) \\ \text{latents} &= f^{-1}(\text{states}; \text{context})\end{aligned}$$

$$q(\text{states} \mid \text{context}) = \frac{\mathcal{N}(\text{latents}; 0, I)}{\left| \det \frac{\partial f}{\partial \text{latents}} \right|}$$


Can **plan** in this distribution!

new idea

Modelling Expert Drivers



optimize



$$\begin{aligned}\text{states} &= f(\text{latents}; \text{context}) \\ \text{latents} &= f^{-1}(\text{states}; \text{context})\end{aligned}$$

$$q(\text{states} \mid \text{context}) = \frac{\mathcal{N}(\text{latents}; 0, I)}{\left| \det \frac{\partial f}{\partial \text{latents}} \right|}$$

Can **plan** in this distribution!

new idea

Planning “Expert-like” Motions to our own Goals

$$\text{planned path} = \arg \max_{\text{states}} \log p(\text{states} \mid \text{goal}, \text{context})$$

Planning “Expert-like” Motions to our own Goals



$$\begin{aligned}\text{planned path} &= \arg \max_{\text{states}} \log p(\text{states} \mid \text{goal}, \text{context}) \\ &= \arg \max_{\text{states}} \log q(\text{states} \mid \text{context}) + \log p(\text{goal} \mid \text{states}, \text{context}) - \log p(\text{goal} \mid \text{context})\end{aligned}$$

Planning “Expert-like” Motions to our own Goals

$$\begin{aligned}\text{planned path} &= \arg \max_{\text{states}} \log p(\text{states} \mid \text{goal}, \text{context}) \\ &= \arg \max_{\text{states}} \log q(\text{states} \mid \text{context}) + \log p(\text{goal} \mid \text{states}, \text{context}) - \log p(\text{goal} \mid \text{context}) \\ &= \arg \max_{\text{states}} \underbrace{\log q(\text{states} \mid \text{context})}_{\text{imitation prior}} + \underbrace{\log p(\text{goal} \mid \text{states}, \text{context})}_{\text{goal likelihood}}\end{aligned}$$

Planning “Expert-like” Motions to our own Goals

$$\begin{aligned}\text{planned path} &= \arg \max_{\text{states}} \log p(\text{states} \mid \text{goal}, \text{context}) \\ &= \arg \max_{\text{states}} \log q(\text{states} \mid \text{context}) + \log p(\text{goal} \mid \text{states}, \text{context}) - \log p(\text{goal} \mid \text{context}) \\ &= \arg \max_{\text{states}} \underbrace{\log q(\text{states} \mid \text{context})}_{\text{imitation prior}} + \underbrace{\log p(\text{goal} \mid \text{states}, \text{context})}_{\text{goal likelihood}}\end{aligned}$$

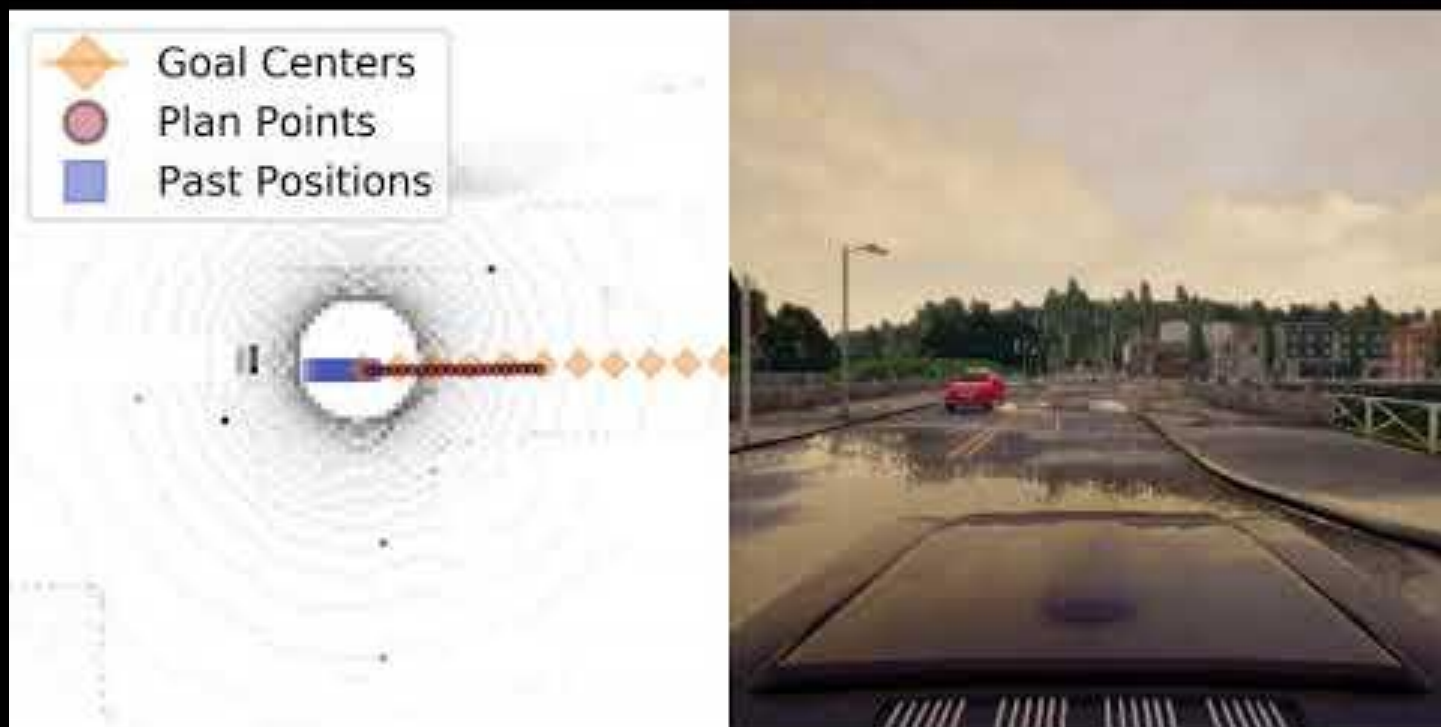
Drive safely...

Planning “Expert-like” Motions to our own Goals

$$\begin{aligned}\text{planned path} &= \arg \max_{\text{states}} \log p(\text{states} \mid \text{goal}, \text{context}) \\ &= \arg \max_{\text{states}} \log q(\text{states} \mid \text{context}) + \log p(\text{goal} \mid \text{states}, \text{context}) - \log p(\text{goal} \mid \text{context}) \\ &= \arg \max_{\text{states}} \underbrace{\log q(\text{states} \mid \text{context})}_{\text{imitation prior}} + \underbrace{\log p(\text{goal} \mid \text{states}, \text{context})}_{\text{goal likelihood}}\end{aligned}$$


Drive safely...

...and get to our destination!









Modelling Multiple Expert Drivers





Modelling Multiple Expert Drivers




Modelling Multiple Expert Drivers

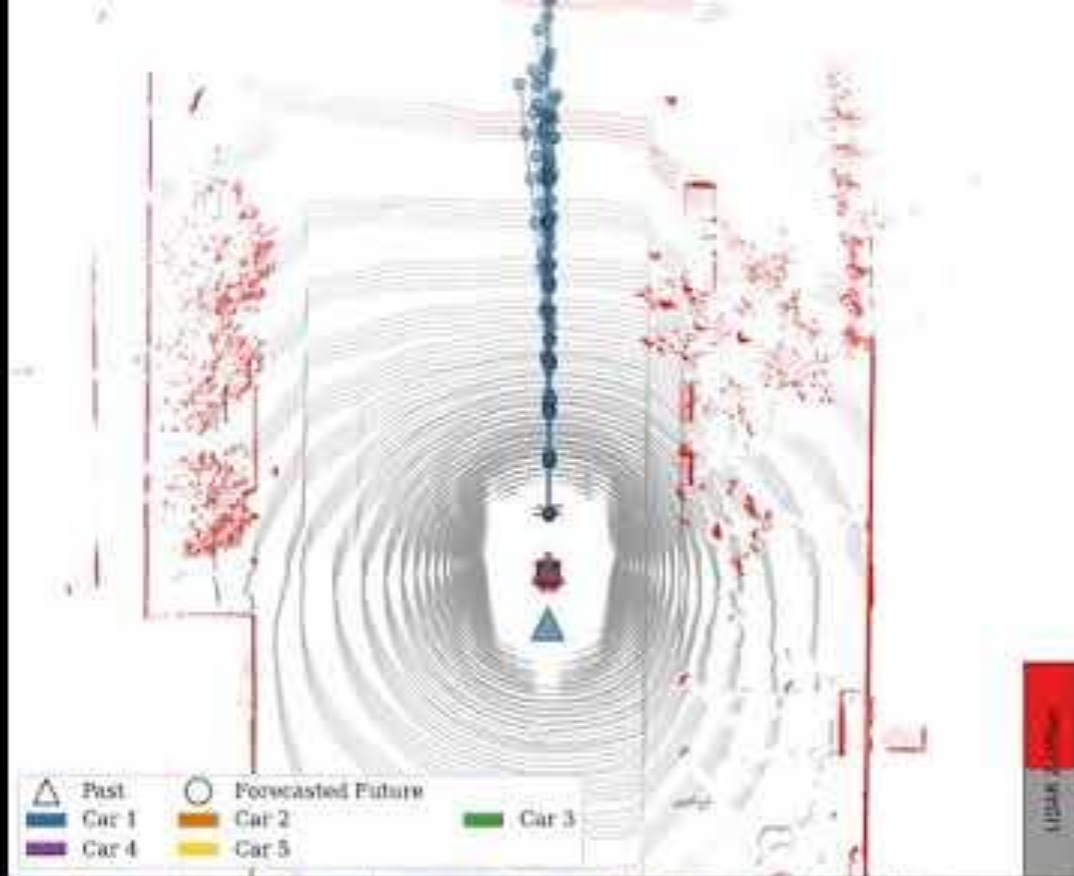


Modelling Multiple Expert Drivers




Forecasting with nuScenes data (Singapore + Boston)






Multi-Agent Planning

Q: but how should ***autonomous vehicles*** predict other agents...**online**?




Multi-Agent Planning

Q: but how should ***autonomous vehicles*** predict other agents...**online**?




Multi-Agent Planning

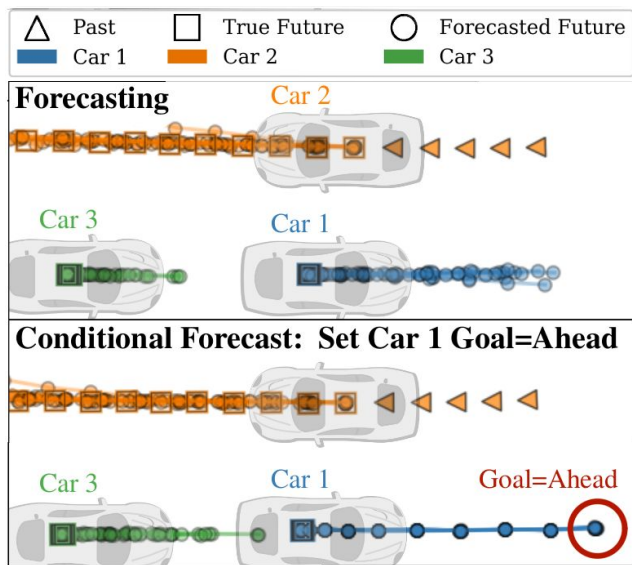
Q: but how should ***autonomous vehicles*** predict other agents...**online**?



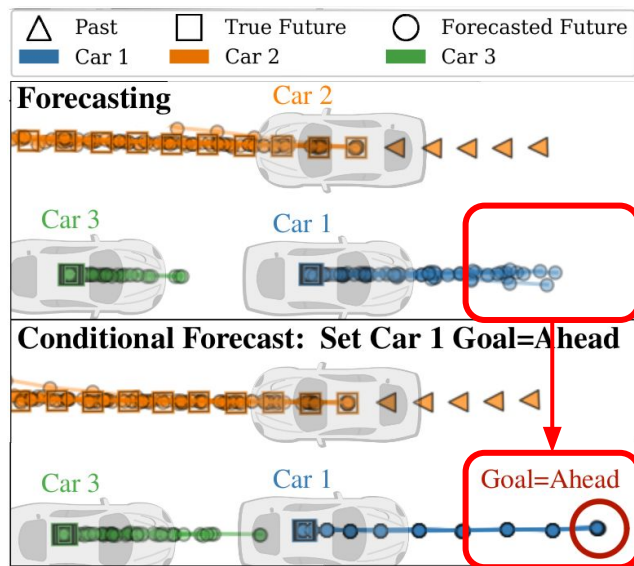
Goal-Conditioned Multi-Agent Forecasting



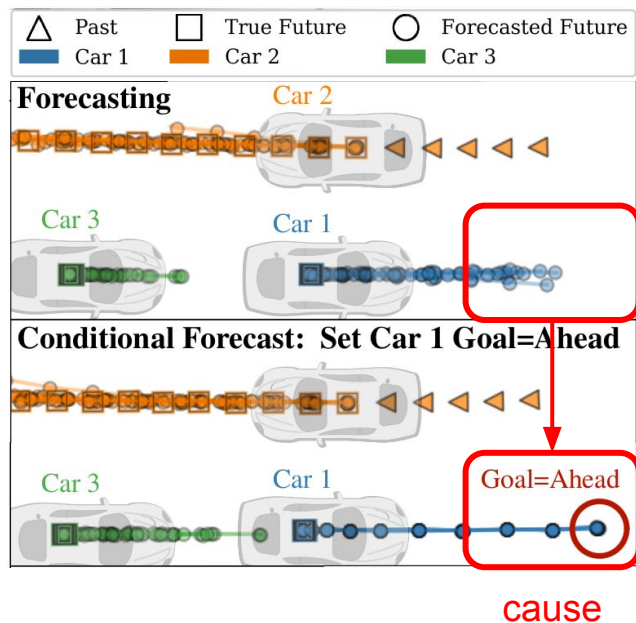
Goal-Conditioned Multi-Agent Forecasting



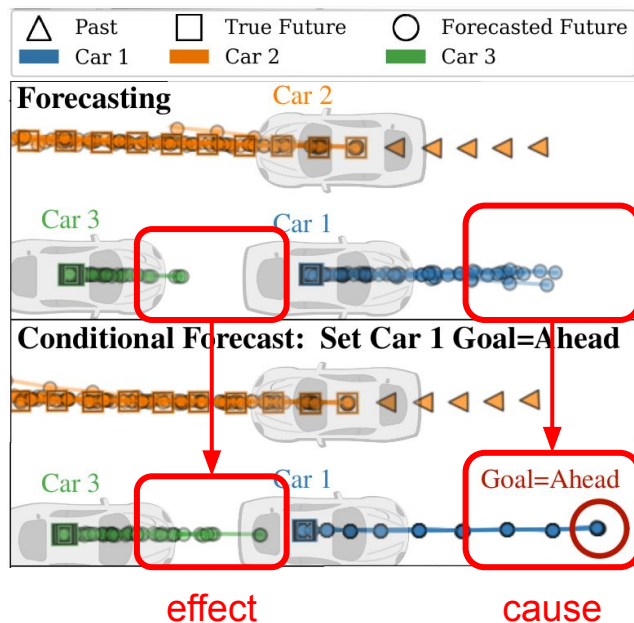
Goal-Conditioned Multi-Agent Forecasting



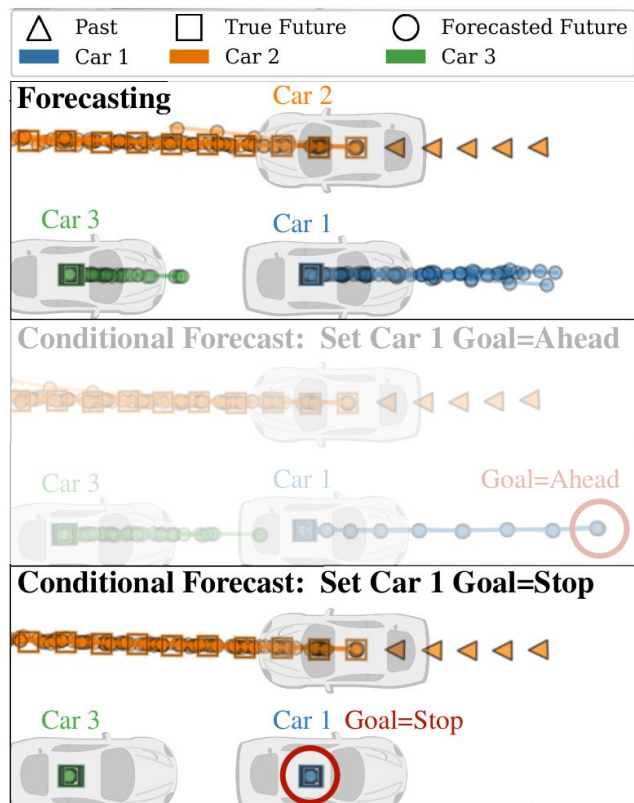
Goal-Conditioned Multi-Agent Forecasting



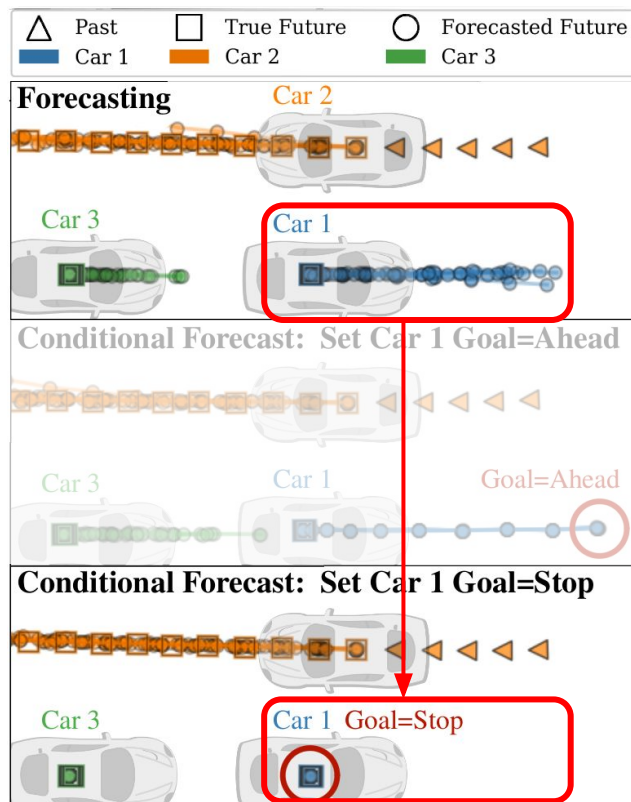
Goal-Conditioned Multi-Agent Forecasting



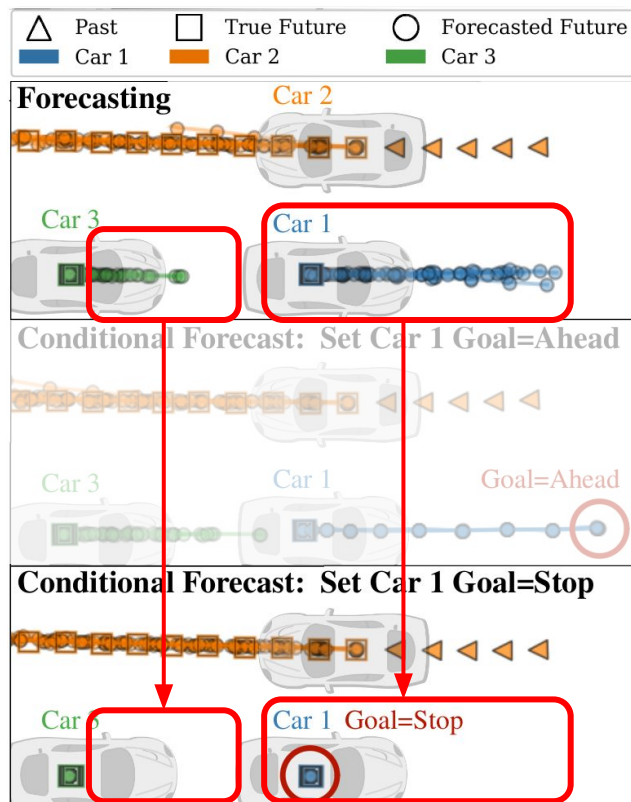
Goal-Conditioned Multi-Agent Forecasting



Goal-Conditioned Multi-Agent Forecasting




Goal-Conditioned Multi-Agent Forecasting




Goal-Conditioned Multi-Agent Forecasting

Forecasting




Conditional forecasting




Goal-Conditioned Multi-Agent Forecasting

Forecasting




Conditional forecasting




How to decide?

Goal-Conditioned Multi-Agent Forecasting

How to decide ?




Goal-Conditioned Multi-Agent Forecasting



$$p \left(\text{states}, \text{states} \mid \text{goal}, \text{context} \right)$$

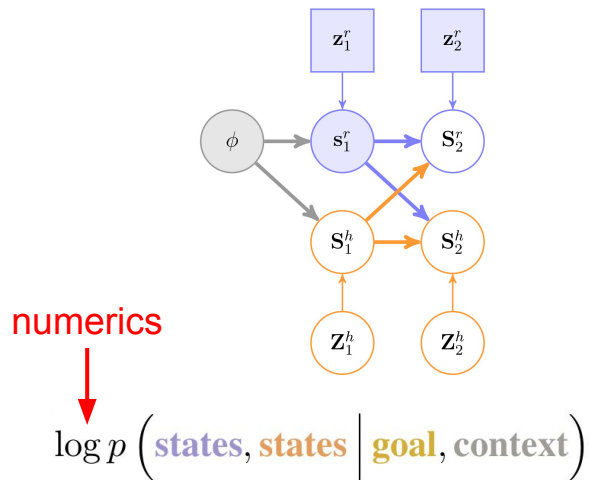
Goal-Conditioned Multi-Agent Forecasting



$$\text{states}, \text{states} = f(\text{latents}, \text{latents}; \text{context})$$


$$p(\text{states}, \text{states} \mid \text{goal}, \text{context})$$

Goal-Conditioned Multi-Agent Forecasting



$$\text{states}, \text{states} = f(\text{latents}, \text{latents}; \text{context})$$

Goal-Conditioned Multi-Agent Forecasting




$$\text{states}, \text{states} = f(\text{latents}, \text{latents}; \text{context})$$

$$\mathbb{E}_{\text{latents}} \left[\log p \left(\text{states}, \text{states} \mid \text{goal}, \text{context} \right) \right]$$

numerics
 ↓
 cannot control
 ↑

Goal-Conditioned Multi-Agent Forecasting



$$\text{states}, \text{states} = f(\text{latents}, \text{latents}; \text{context})$$

$$\arg \max_{\text{latents}} \mathbb{E}_{\text{latents}} \left[\log p \left(\text{states}, \text{states} \mid \text{goal}, \text{context} \right) \right]$$


numerics (points to the log p term)

cannot control (points to the $\mathbb{E}_{\text{latents}}$ term)

can control (points to the $\arg \max_{\text{latents}}$ term)

Goal-Conditioned Multi-Agent Forecasting


How we decide



$$\text{states}, \text{states} = f(\text{latents}, \text{latents}; \text{context})$$

$$\text{latents}^* = \arg \max_{\text{latents}} \mathbb{E}_{\text{latents}} \left[\log p(\text{states}, \text{states} \mid \text{goal}, \text{context}) \right]$$


Goal-Conditioned Multi-Agent Forecasting



$$\text{states}, \text{states} = f(\text{latents}, \text{latents}; \text{context})$$

$$\begin{aligned} \text{latents}^* &= \arg \max_{\text{latents}} \mathbb{E}_{\text{latents}} \left[\log p \left(\text{states}, \text{states} \mid \text{goal}, \text{context} \right) \right] \\ &= \arg \max_{\text{latents}} \mathbb{E}_{\text{latents}} \left[\log q \left(\text{states}, \text{states} \mid \text{context} \right) + \log p(\text{goal} \mid \text{states}, \text{states}, \text{context}) - \log p(\text{goal} \mid \text{context}) \right] \end{aligned}$$

Goal-Conditioned Multi-Agent Forecasting



$$\text{states}, \text{states} = f(\text{latents}, \text{latents}; \text{context})$$

$$\begin{aligned} \text{latents}^* &= \arg \max_{\text{latents}} \mathbb{E}_{\text{latents}} \left[\log p \left(\text{states}, \text{states} \mid \text{goal}, \text{context} \right) \right] \\ &= \arg \max_{\text{latents}} \mathbb{E}_{\text{latents}} \left[\log q \left(\text{states}, \text{states} \mid \text{context} \right) + \log p \left(\text{goal} \mid \text{states}, \text{states}, \text{context} \right) - \log p \left(\text{goal} \mid \text{context} \right) \right] \\ &= \arg \max_{\text{latents}} \mathbb{E}_{\text{latents}} \left[\underbrace{\log q \left(\text{states}, \text{states} \mid \text{context} \right)}_{\text{multi-agent prior}} + \underbrace{\log p \left(\text{goal} \mid \text{states}, \text{states}, \text{context} \right)}_{\text{goal likelihood}} \right] \end{aligned}$$

Future work: how to respond to “out-of-distribution” scenes?



Thank you!

Single-agent forecasting + control

Deep Imitative Models for Flexible Inference, Planning, and Control

Nicholas Rhinehart, Rowan McAllister, Sergey Levine

<http://imitate.ml>

Multi-agent forecasting

PRECOC: PREDiction Conditioned On Goals in Visual Multi-Agent Settings


Nicholas Rhinehart, Rowan McAllister, Kris Kitani, Sergey Levine

<http://precog.ml>


Bayesian single-agent forecasting + control

(in progress)


Forecasting with nuScenes data (Singapore + Boston)



Car 2 is predicted to overtake **Car 1**, which itself is forecasted to continue to wait for pedestrians and **Car 2**.




Car 4 is predicted to wait for a clear intersection, and **Car 5** is predicted to either start turning or continue straight.



Forecasting

PRECOG




Forecasting

PRECOG

Goal-Conditioned Forecasting

Forecasting



Car 1 Forecasts Stopping

